

# Attention Mechanism and Transformers

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Agenda

- Transformers Quiz
- Introduction to Transformers
- Positional Embedding
- Multi-head Attention
- Masking
- Encoder-Decoder Attention
- Transformer Architectures

**Let's begin the discussion by answering a few questions  
on Attention mechanism and Transformers**

# Transformers Quiz

**Considering the transformer architecture proposed In the original paper (Attention Is All You Need, 2017), which of the following statements are true?**

**A**

A transformer is a neural network

**B**

The encoder stage outputs a latent representation of the input

**C**

The decoder stage outputs a sequence of tokens based on the latent representation from the encoder stage

**D**

The encoder and decoder stages always consist of one encoder and decoder block only

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Transformers Quiz

**Considering the transformer architecture proposed In the original paper (Attention Is All You Need, 2017), which of the following statements are true?**

A

A transformer is a neural network

B

The encoder stage outputs a latent representation of the input

C

The decoder stage outputs a sequence of tokens based on the latent representation from the encoder stage

D

The encoder and decoder stages always consist of one encoder and decoder block only

This file is meant for personal use by erikvdesigner@gmail.com only.

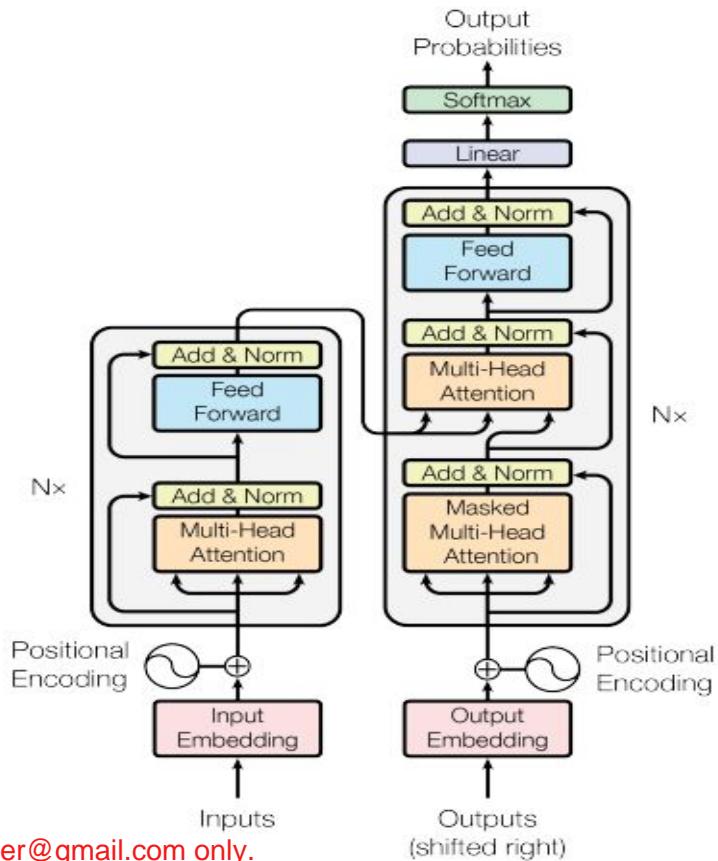
Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Transformers

Transformers are a **type of neural network architecture**

Transformers were **introduced** in a paper by **Vaswani et al. in 2017**

Transformers are based on the idea of **self-attention**



**Source:** Image from the original research paper [Attention Is All You Need](#)

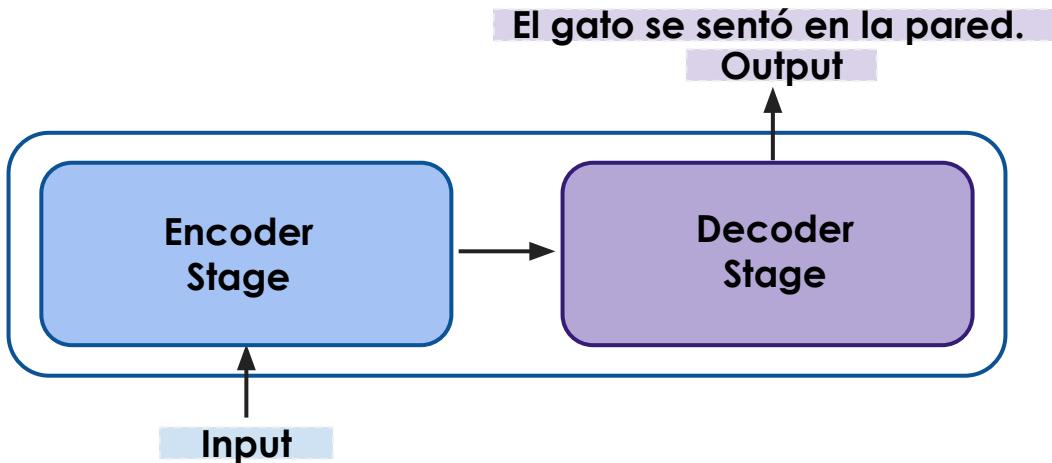
This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Transformers - Working

The **encoder** takes in a sequence of tokens (e.g. words or characters) and outputs a **latent representation**

The **decoder** then takes this latent representation as input and outputs a **sequence of tokens**

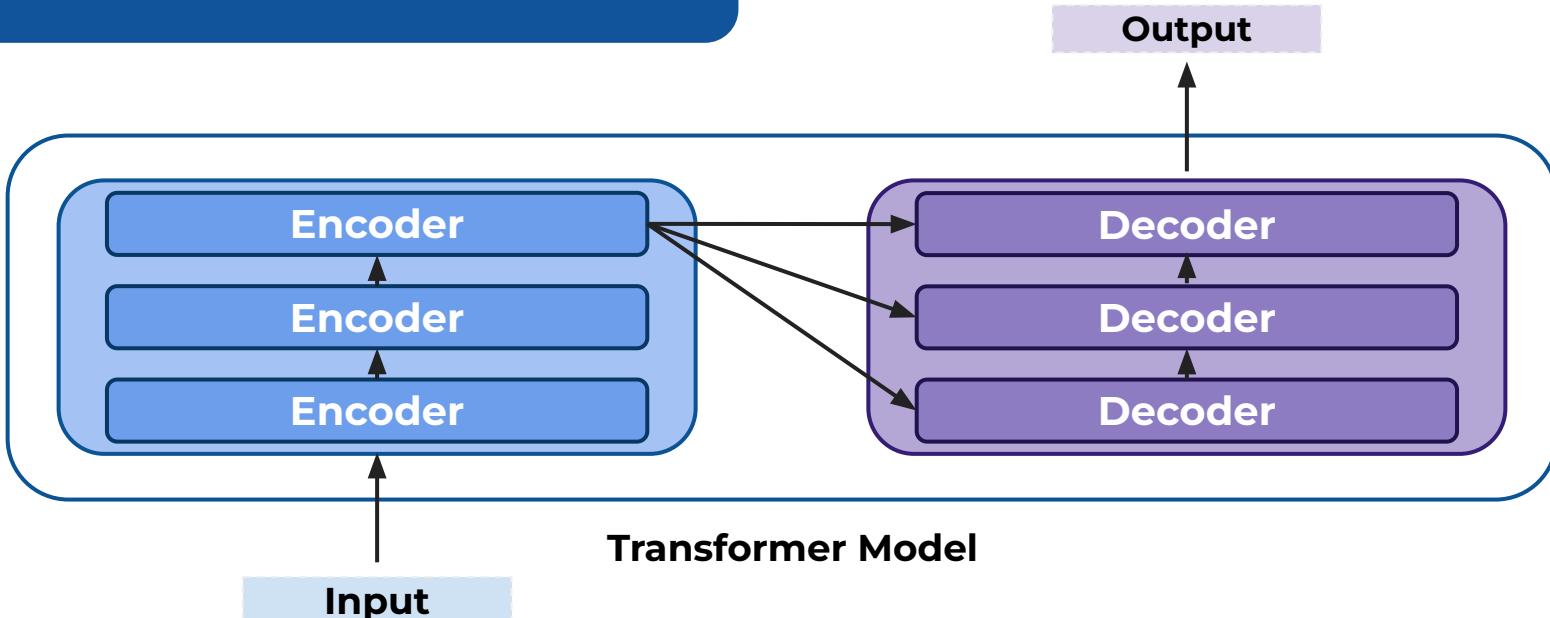


This content is for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Transformers - Working

In reality, the **Encoder** and **Decoder** stage each comprise of **several individual blocks** of Encoders and Decoders.



# Transformers Quiz

**Which of the following best describes the purpose of positional encoding in transformer ?**

A

It introduces randomness to word embeddings to enhance model generalization.

B

It helps the model differentiate between words with similar semantic meanings.

C

It provides the model with information about the order of words in a sequence.

D

It reduces the computational complexity of the self-attention mechanism.

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Transformers Quiz

**Which of the following best describes the purpose of positional encoding in transformer ?**

A

It introduces randomness to word embeddings to enhance model generalization.

B

It helps the model differentiate between words with similar semantic meanings.

C

It provides the model with information about the order of words in a sequence.

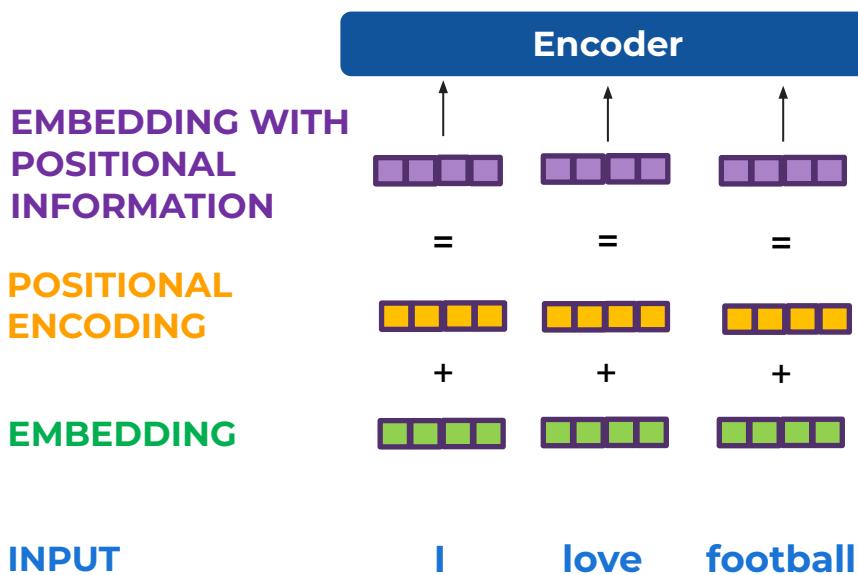
D

It reduces the computational complexity of the self-attention mechanism.

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Positional Encoding



**Positional Encoding** is a way to account for the order of the words in the input sequence.

**Positional Encoding** is a vector added to each input embedding.

These vectors follow a specific pattern that the model learns, which helps it determine the position of each word, or the distance between different words in the sequence.

# Transformers Quiz

**Which of the following best describes the purpose of self-attention mechanism in transformer?**

A

It helps the model focus on relevant parts of the input sequence when making predictions.

B

It reduces the computational complexity of the neural network.

C

It enables the model to generate synthetic data for training purposes.

D

It increases the model's ability to generalize to unseen data.

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Transformers Quiz

**Which of the following best describes the purpose of self-attention mechanism in transformer?**

A

It helps the model focus on relevant parts of the input sequence when making predictions.

B

It reduces the computational complexity of the neural network.

C

It enables the model to generate synthetic data for training purposes.

D

It increases the model's ability to generalize to unseen data.

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Self-Attention Mechanism

The **self-attention mechanism** lies at the **core** of **transformer models**

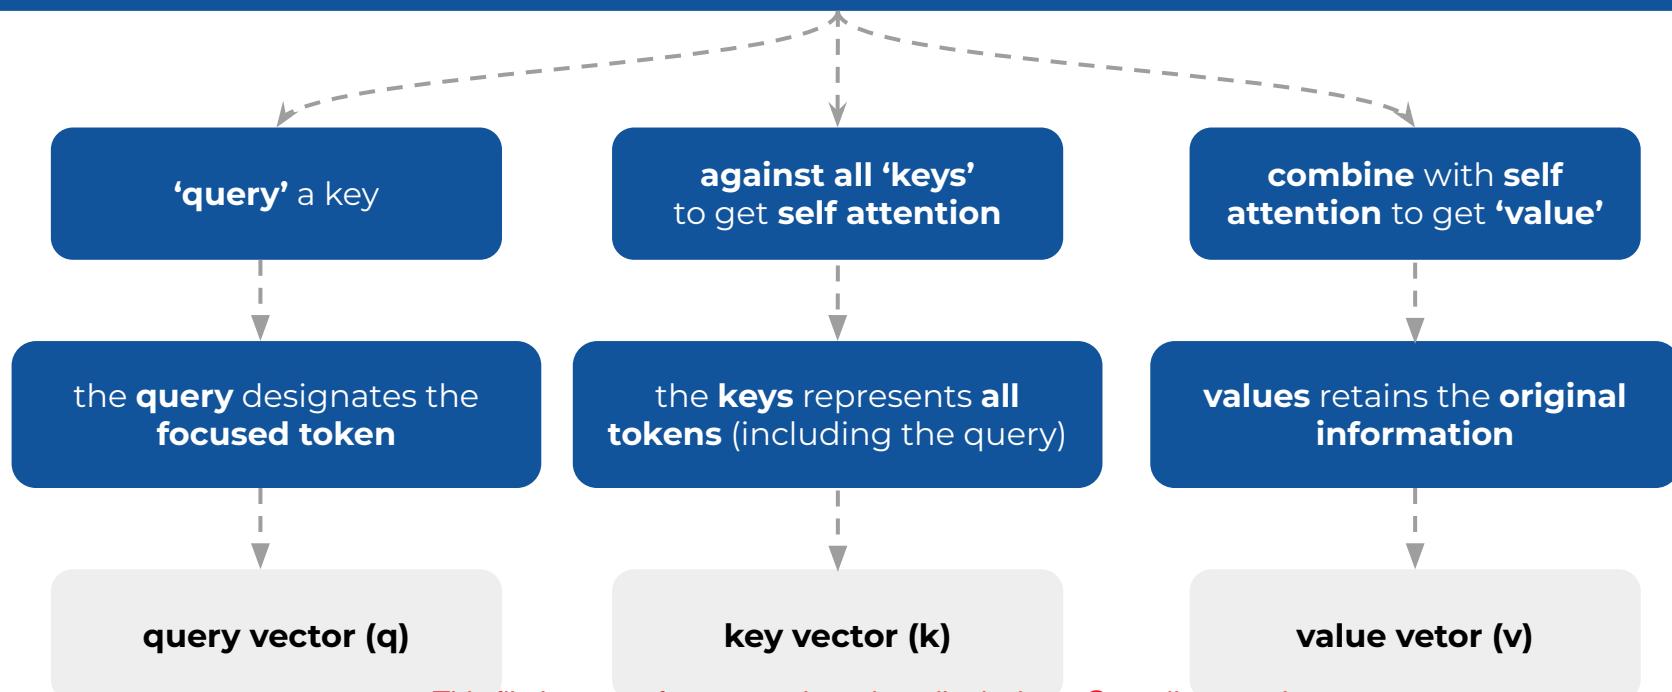
Self attention allows us to generate a **context-aware representation** of **each token** in the input

The **context-aware representation** of **each token** is generated with respect to all other tokens in the input

The context-aware representation **focuses** on the **relevant parts of the input** for a given task

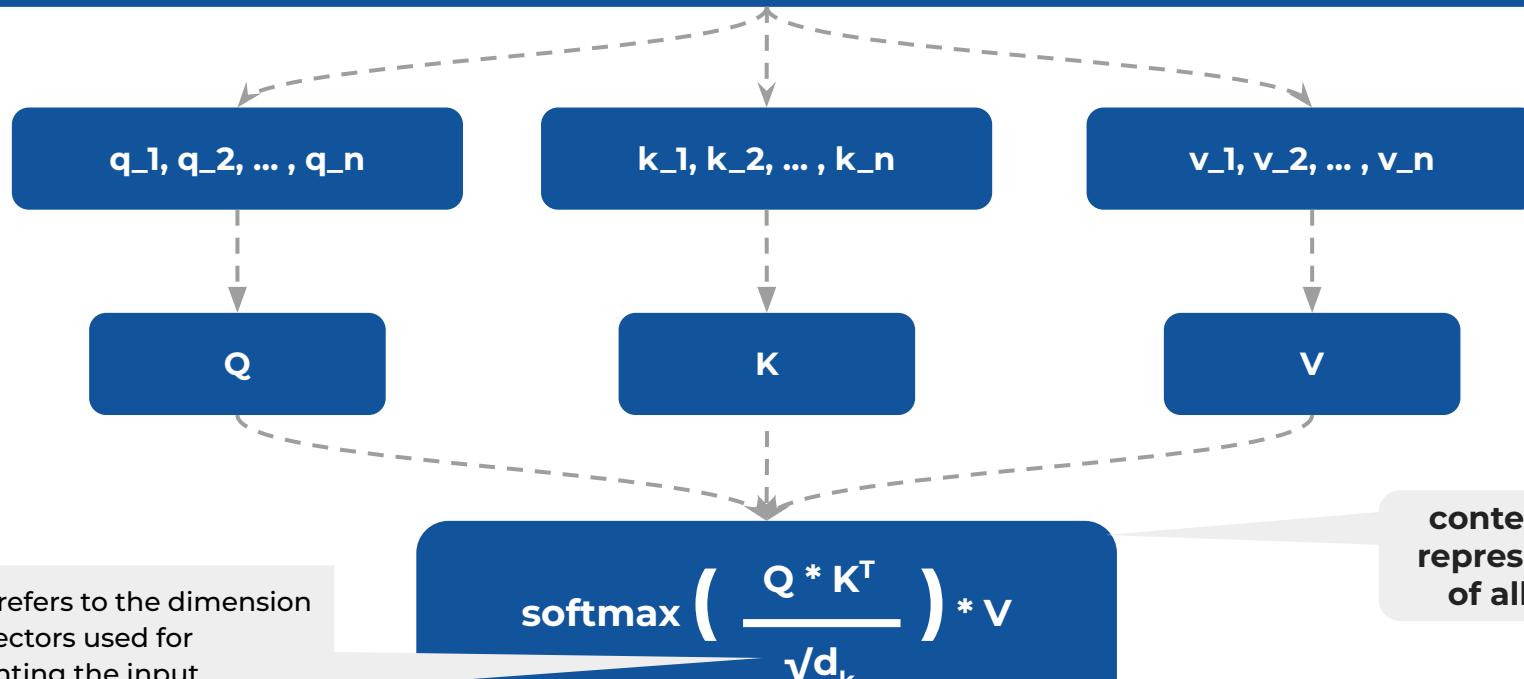
# Self-Attention - Computation

The **steps** to get the **context-aware representation** for **each** of the **token** in the **input** are



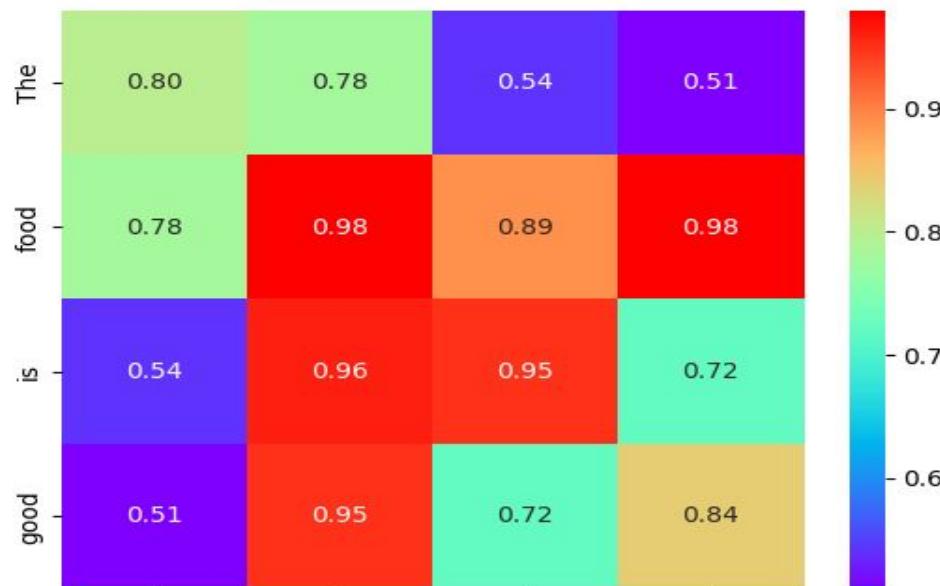
# Self-Attention - Computation

We **stack** these query, key, and value **vectors** into **matrices**



## Self-Attention - Example

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$



# Transformers Quiz

In a transformer model with multi-head attention, if there are 8 attention heads and each head has a dimensionality of 64, what is the dimension of the multi-head attention output after concatenation?

A

64

B

128

C

256

D

512

# Transformers Quiz

In a transformer model with multi-head attention, if there are 8 attention heads and each head has a dimensionality of 64, what is the dimension of the multi-head attention output after concatenation?

A

64

B

128

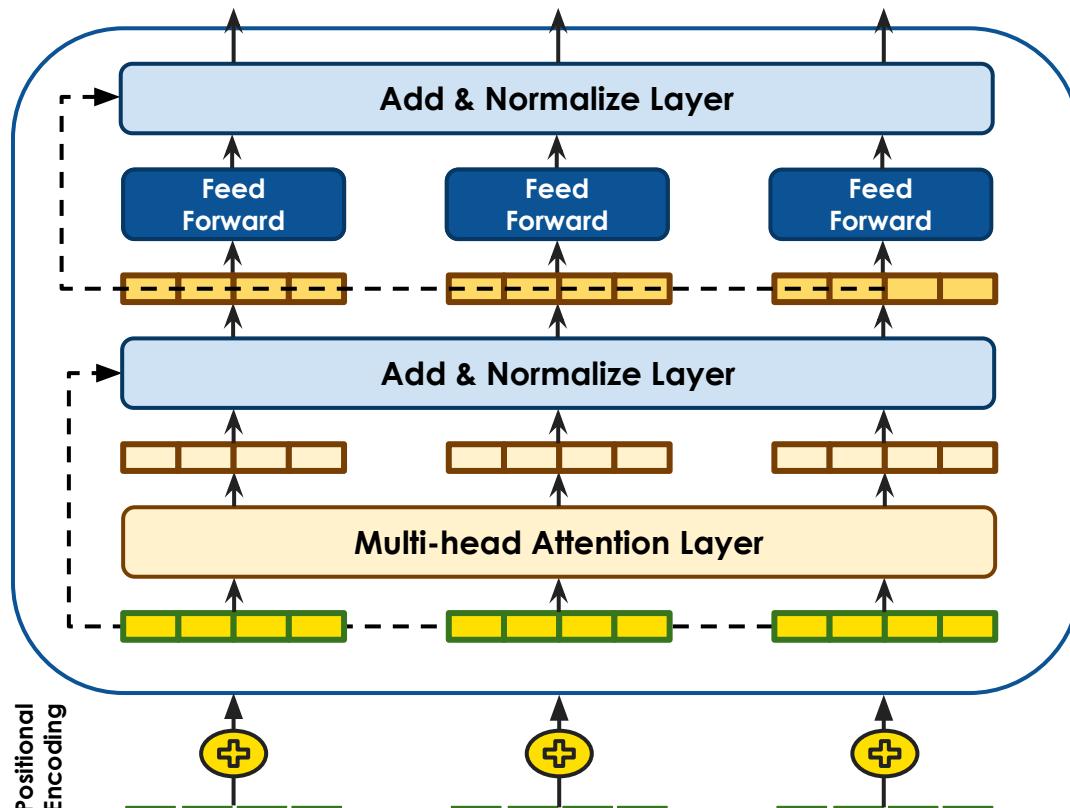
C

256

D

512

# Transformer Architecture - Encoder Block



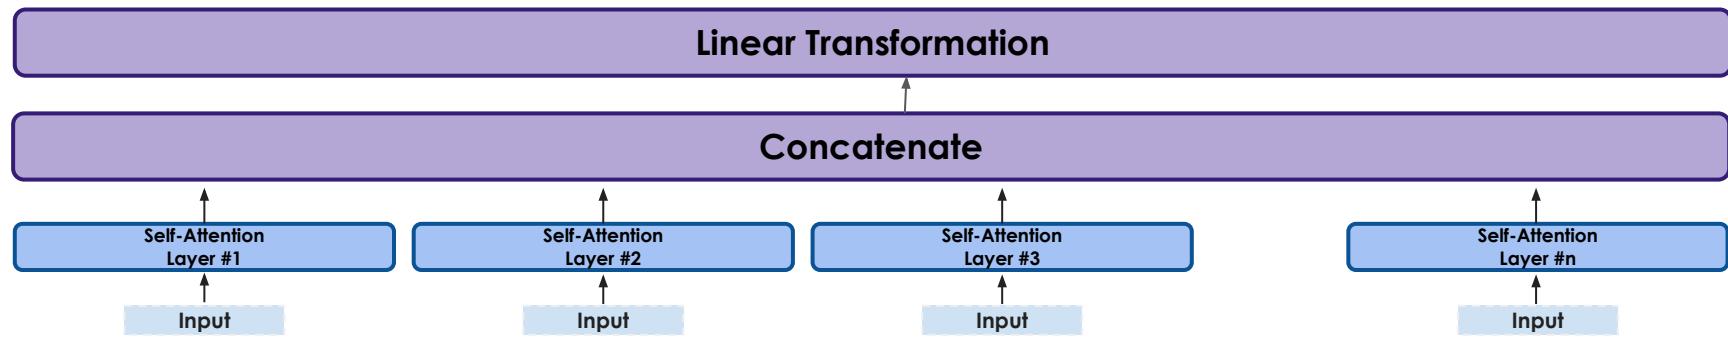
This file is meant for personal use by erkvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Multi-Head Attention

The output of each self-attention layer is taken and concatenated

The linear transformation layer is merely a fully-connected layer of neurons



Each head produces a 64 dimensional vector. Since there are 8 such heads, after concatenation, it will result into a  $8 \times 64 = 512$  dimensional vector.

# Transformers Quiz

**Why is masking necessary in the decoder of a Transformer model for sequence-to-sequence tasks?**

A

To limit the decoder's access to information from the future positions

B

To enhance the attention mechanisms focus on relevant information

C

To prevent overfitting during training

D

To increase the model's capacity to handle longer sequences

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Transformers Quiz

**Why is masking necessary in the decoder of a Transformer model for sequence-to-sequence tasks?**

A

To limit the decoder's access to information from the future positions

B

To enhance the attention mechanisms focus on relevant information

C

To prevent overfitting during training

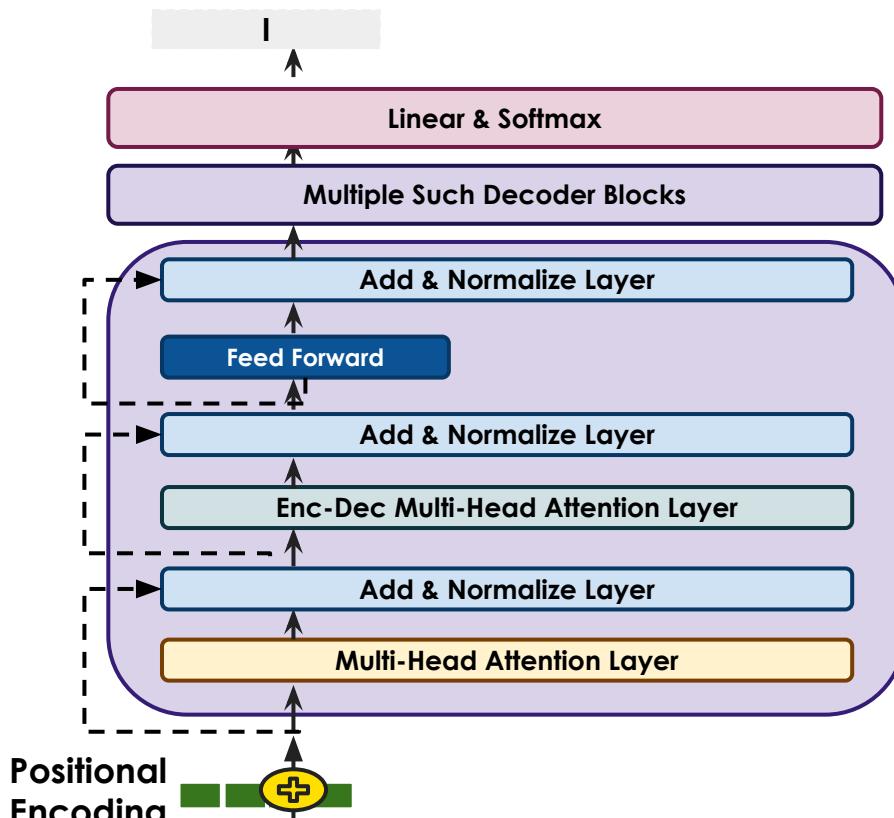
D

To increase the model's capacity to handle longer sequences

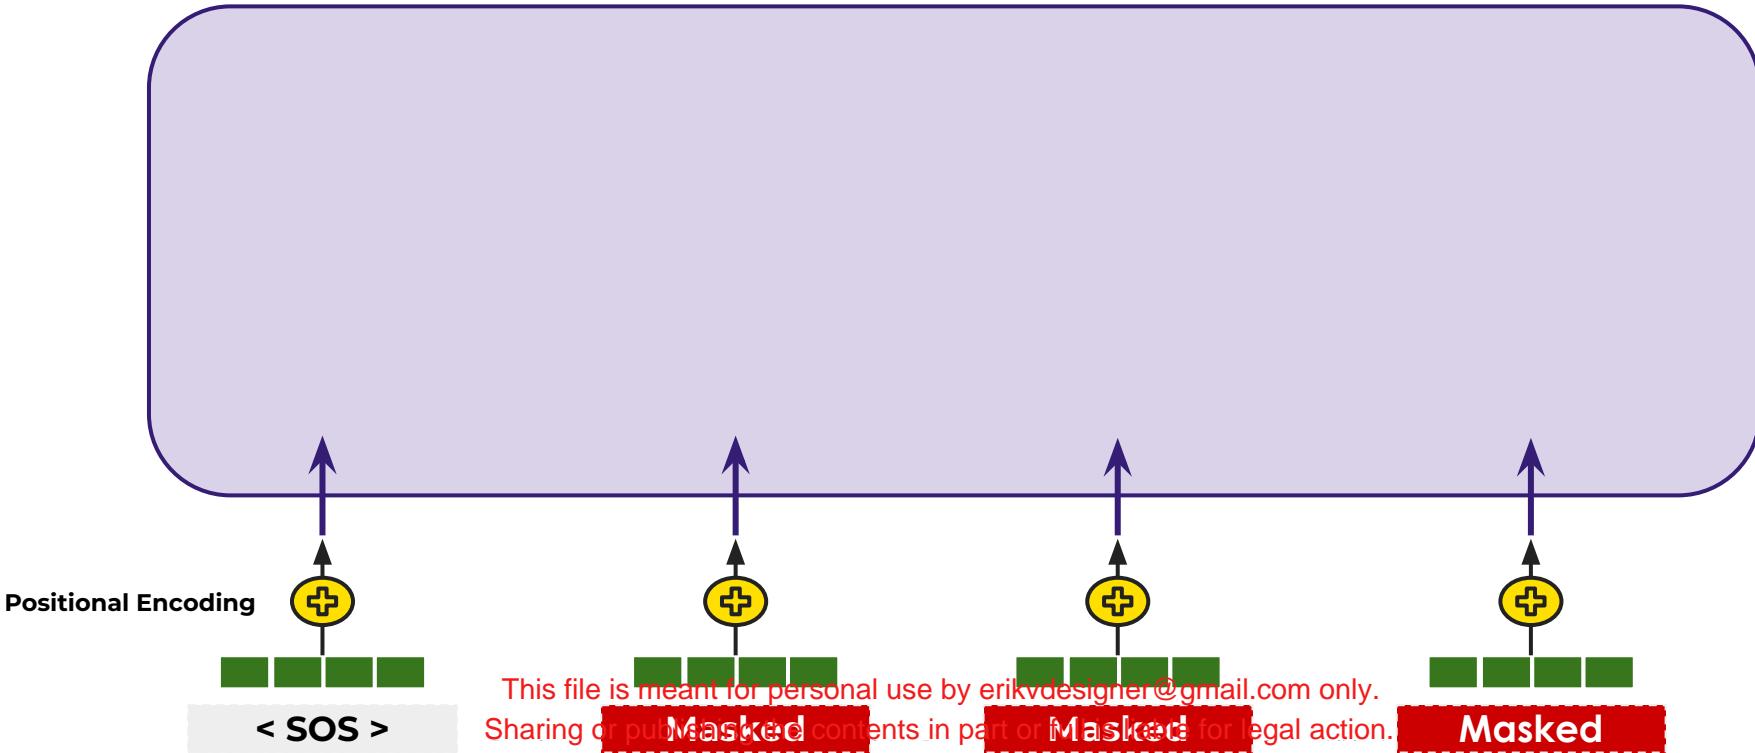
This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

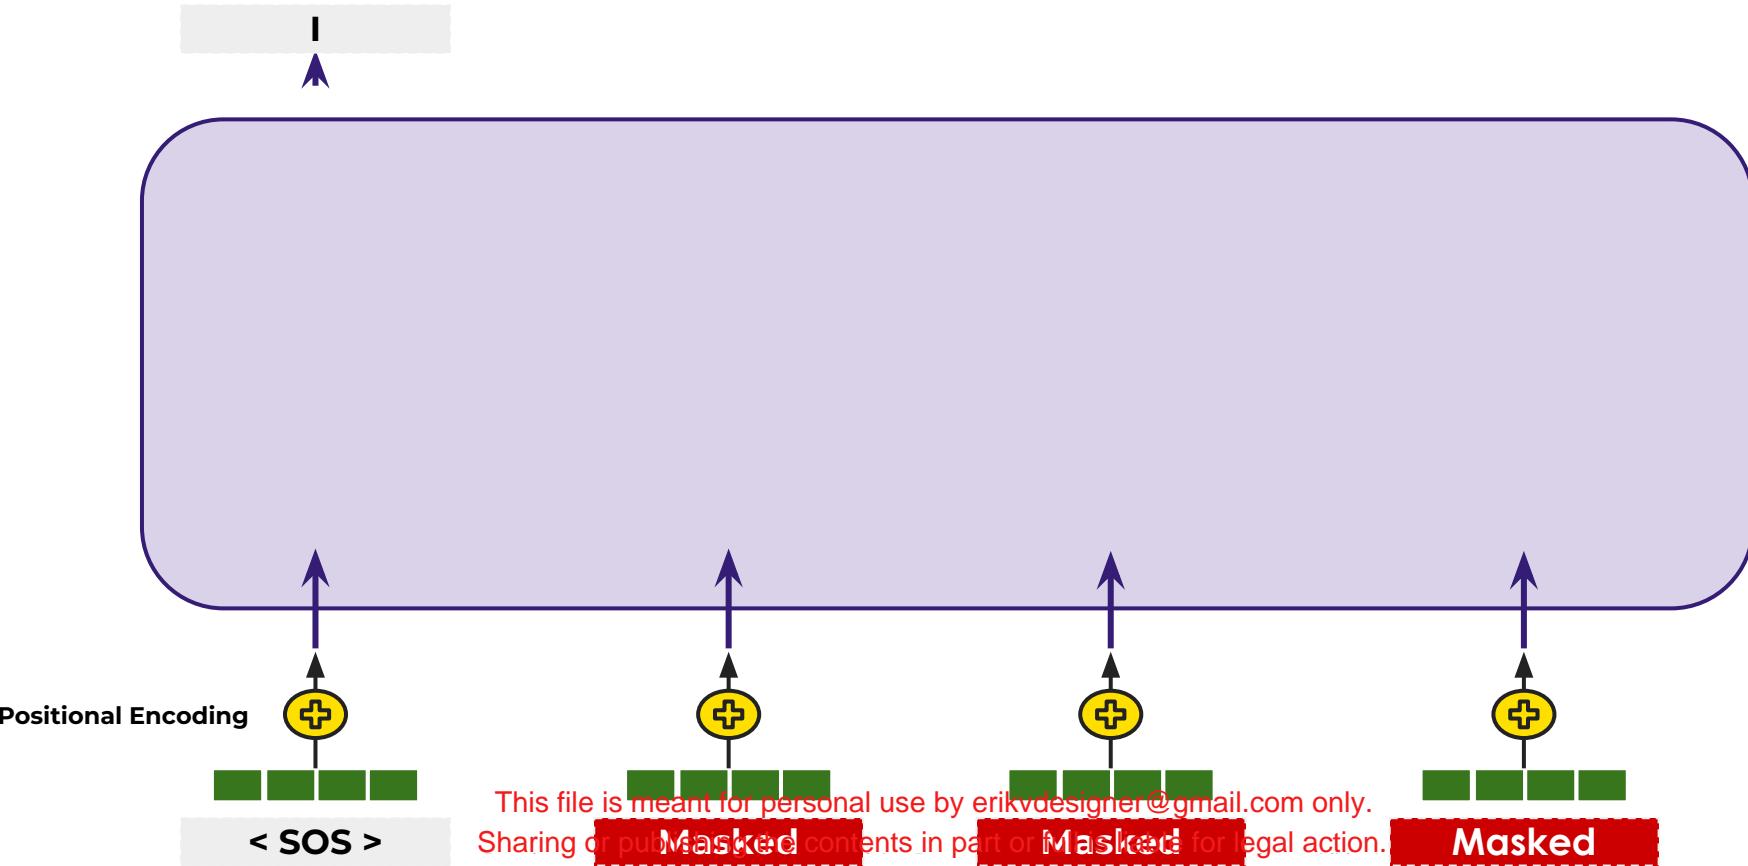
# Transformer Architecture - Decoder Block



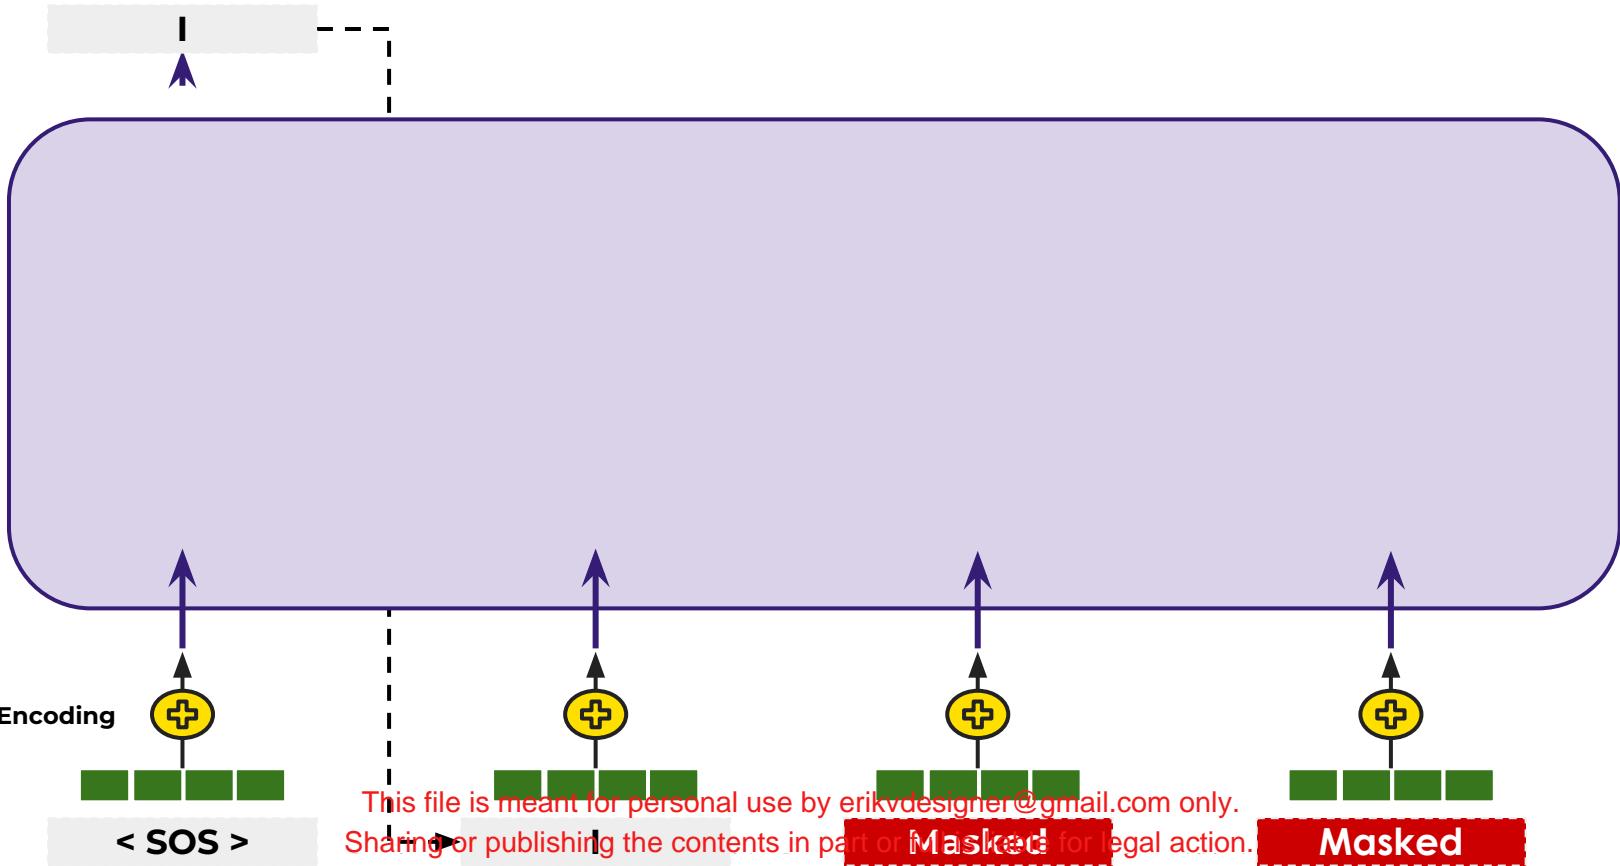
# Decoder - Masking



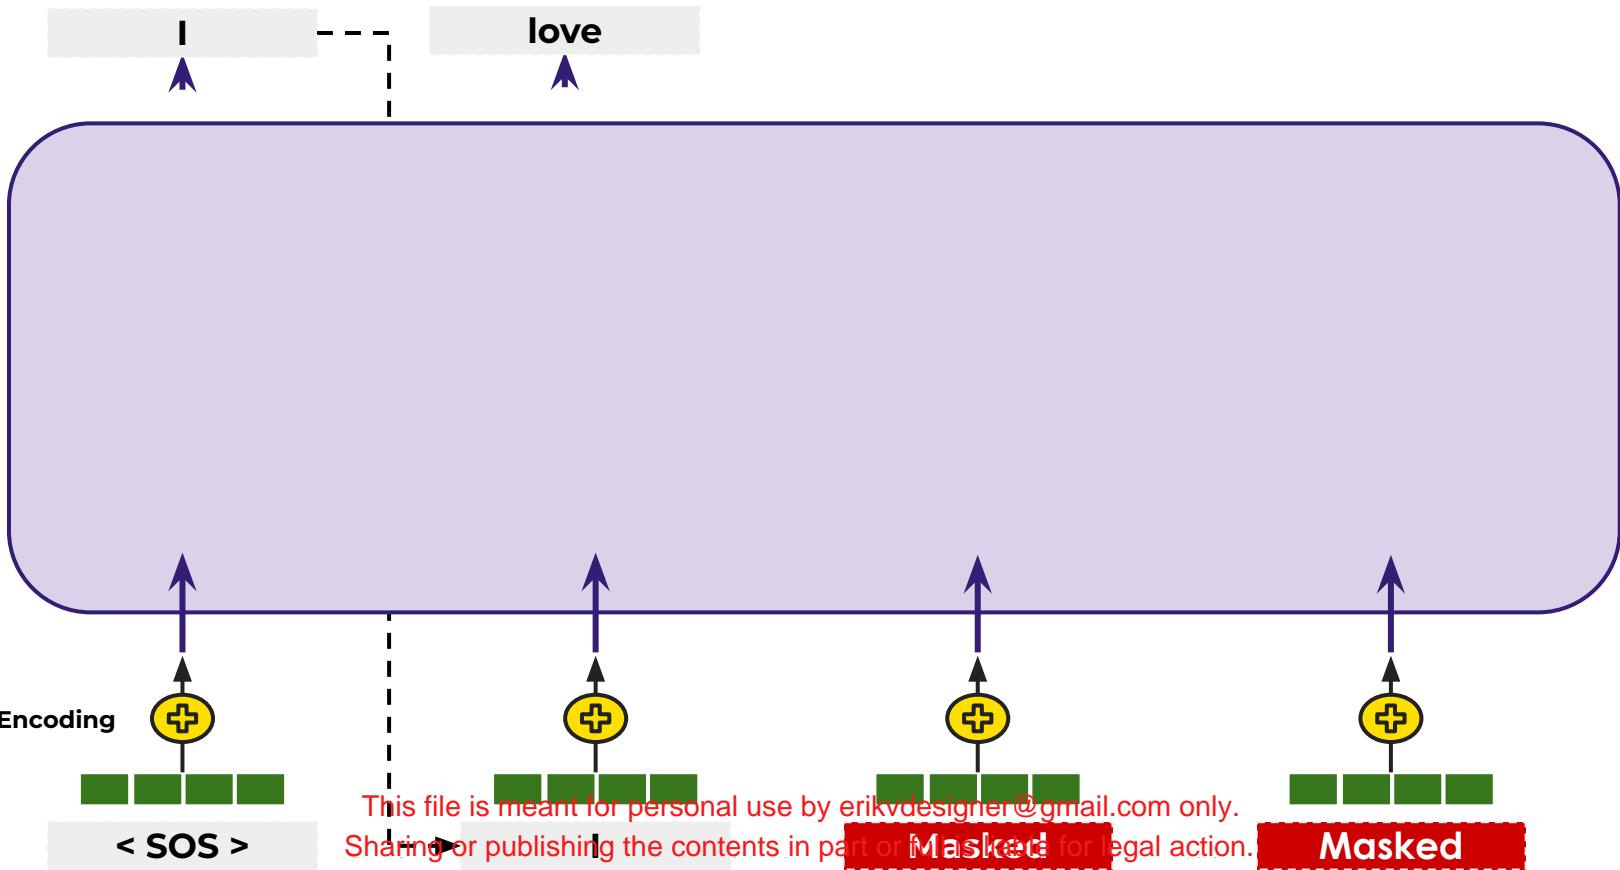
# Decoder - Masking



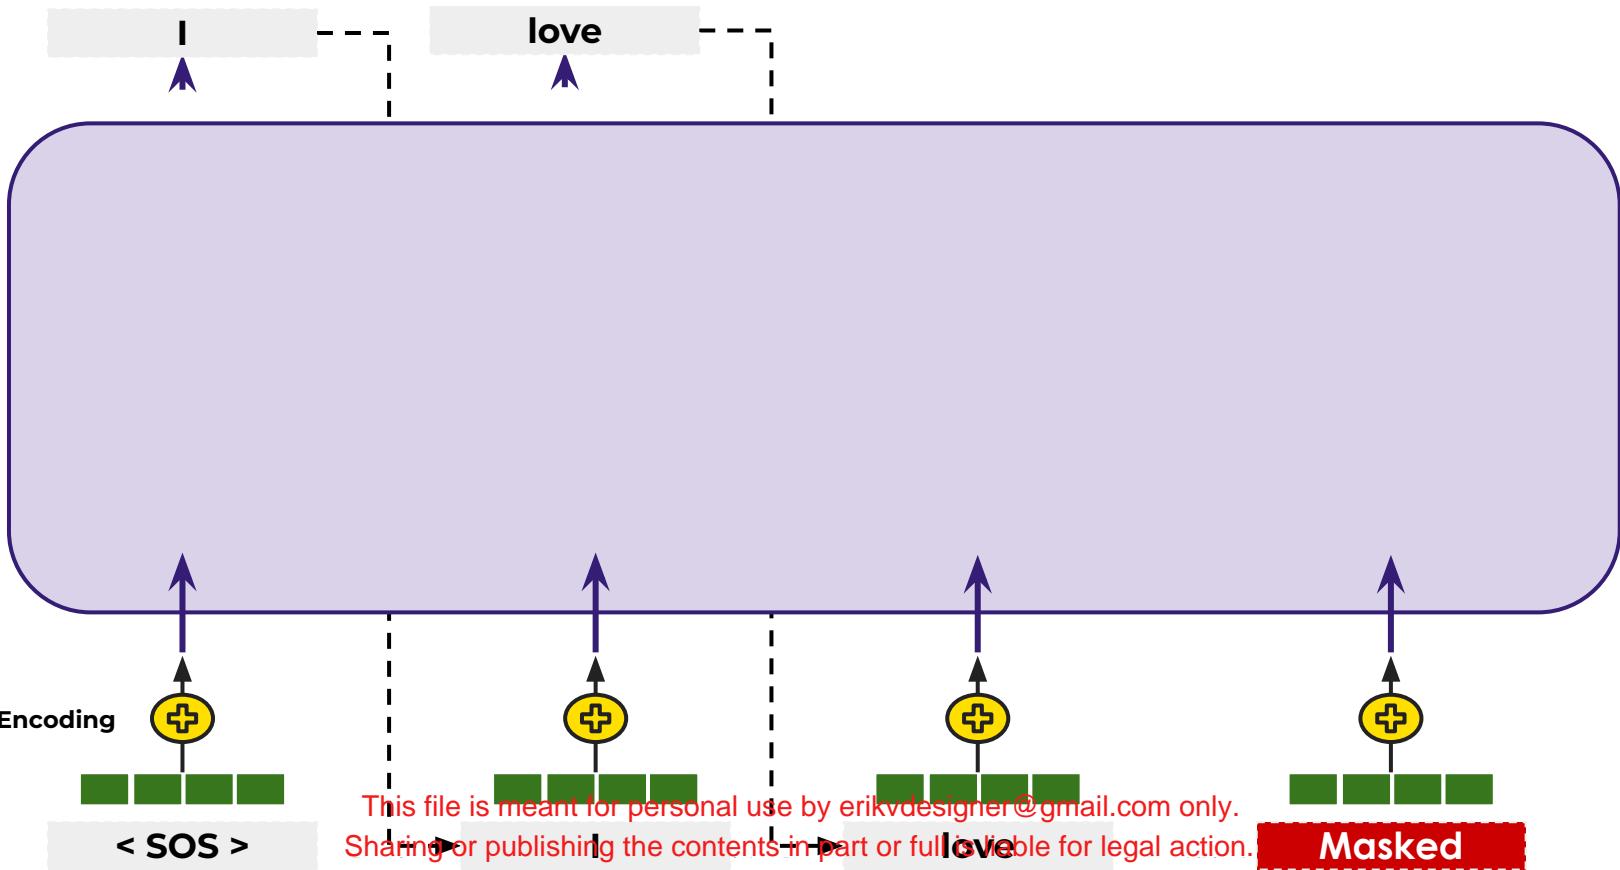
## Decoder - Masking



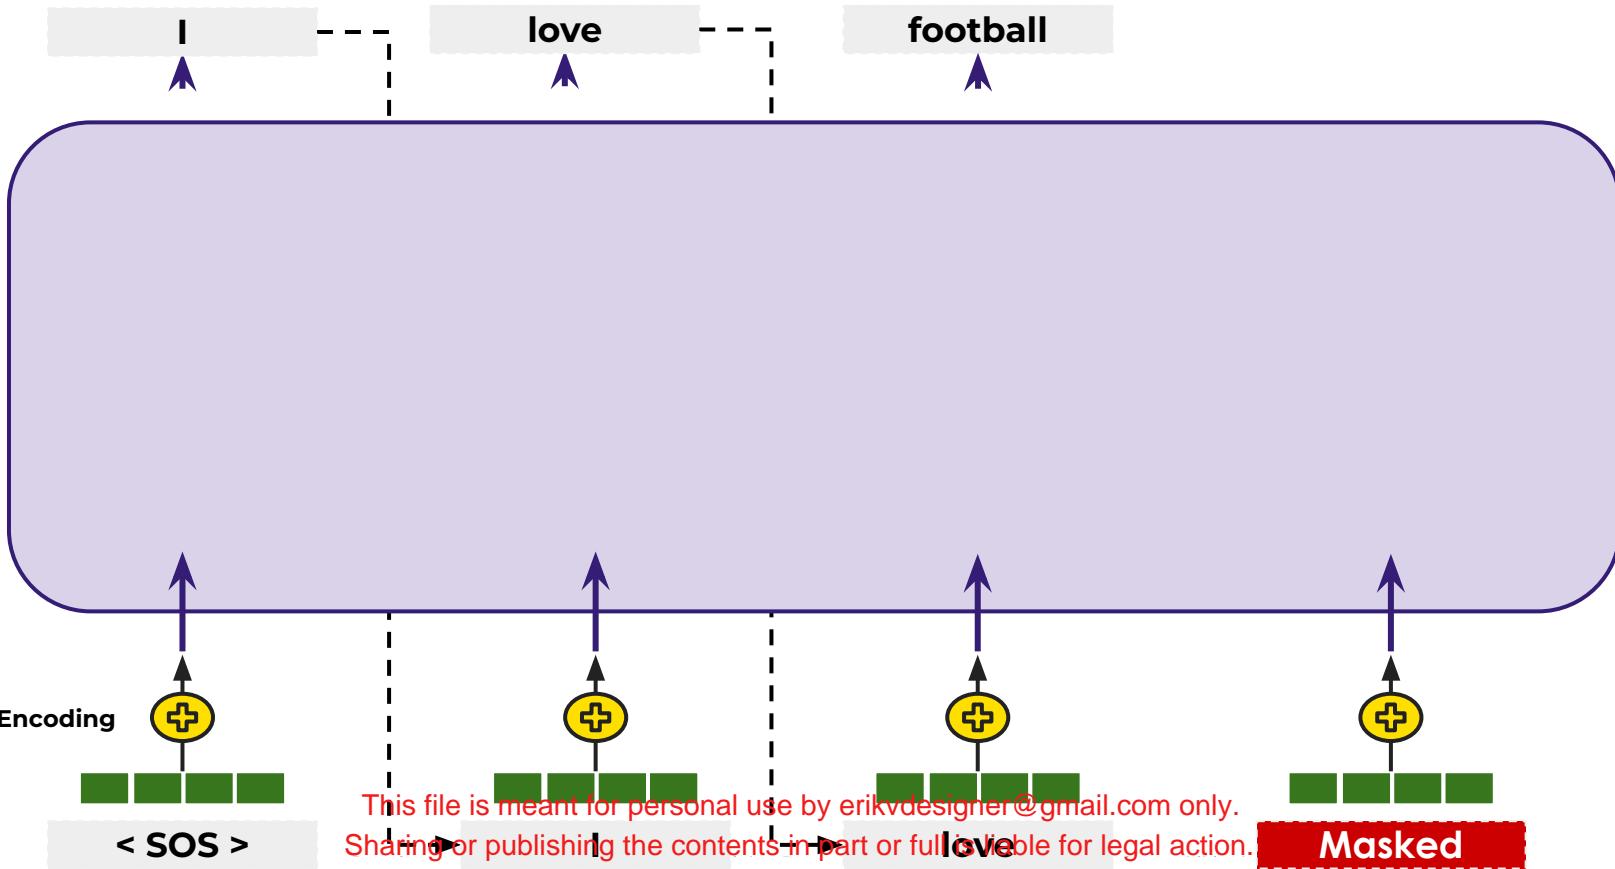
## Decoder - Masking



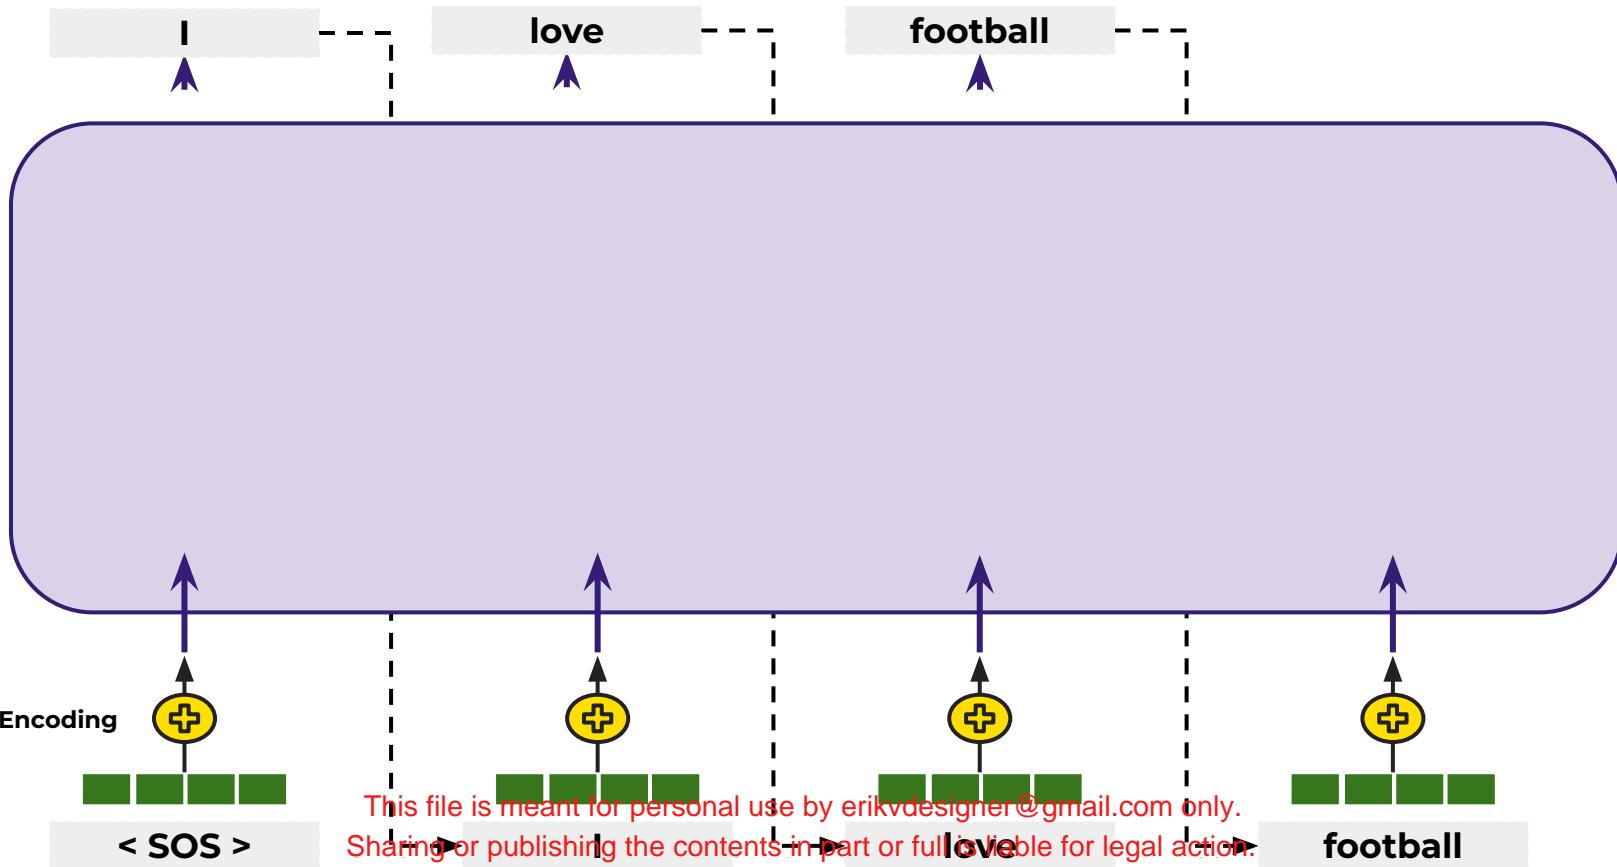
## Decoder - Masking



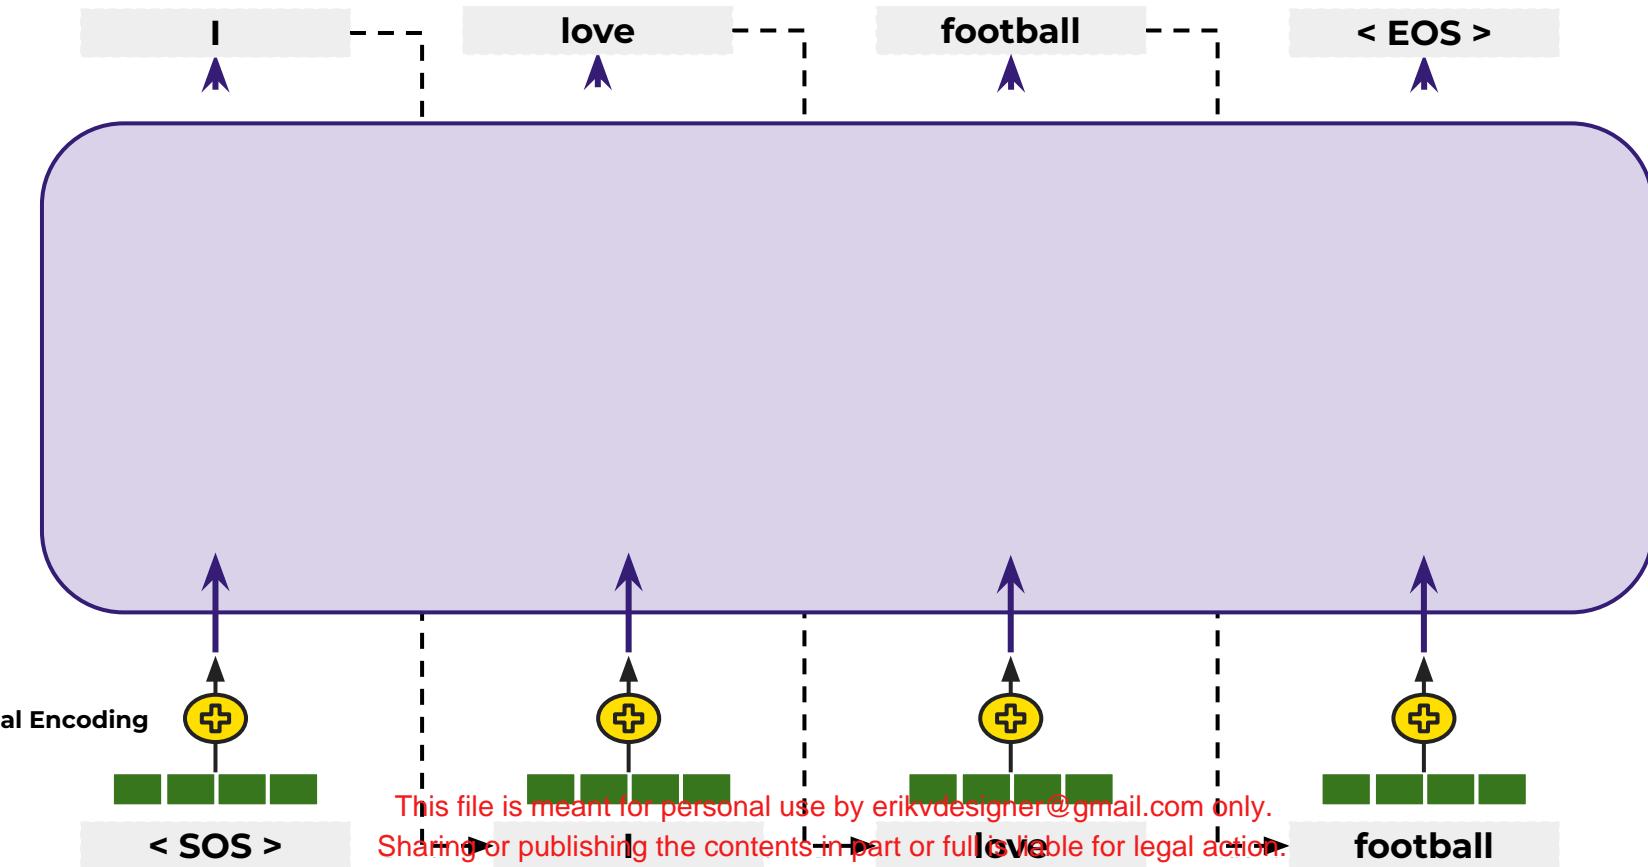
## Decoder - Masking



## Decoder - Masking



## Decoder - Masking

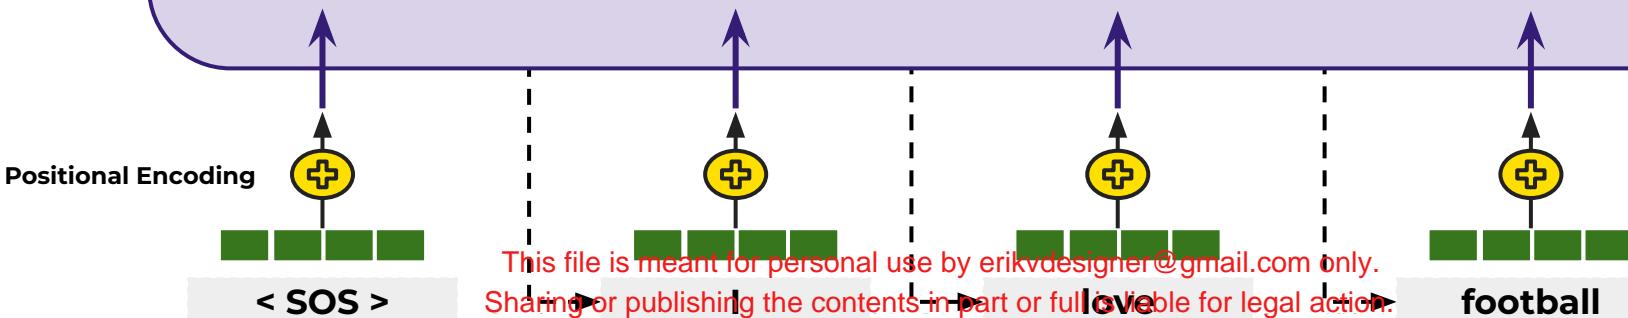


# Decoder - Masking



This characteristic of “**masking**” the future words / tokens and only allowing inputs to the Decoder operations from current & past words in each run through the Decoder, is why this process is sometimes called **Masked Self-Attention**.

**Note:** For each time step, **not just the input from that word, but the inputs of all previous words also go into the decoder**, to predict the output of that timestep.



# Transformers Quiz

**Which of the following accurately describes the origin of the query, key, and value in encoder-decoder attention mechanisms?**

A

The query, key, and value are all derived solely from the encoder

B

The query is derived from the decoder, while the key and value are from the encoder

C

The query and value are derived from the encoder, while the key is from the decoder

D

The query, key, and value are all derived from the decoder

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# Transformers Quiz

**Which of the following accurately describes the origin of the query, key, and value in encoder-decoder attention mechanisms?**

A

The query, key, and value are all derived solely from the encoder

B

The query is derived from the decoder, while the key and value are from the encoder

C

The query and value are derived from the encoder, while the key is from the decoder

D

The query, key, and value are all derived from the decoder

This file is meant for personal use by erikvdesigner@gmail.com only.

Sharing or publishing the contents in part or full is liable for legal action.  
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

# The Encoder-Decoder Attention Layer

The difference from normal Self-Attention is that in this layer, **the K and V vectors are not generated from the input embeddings to this layer**, the way they were in the normal Self-Attention layer.



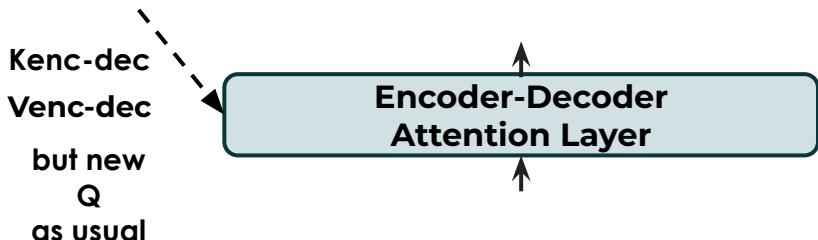
In fact, we utilize a **K encoder-decoder (K enc-dec)** and a **V encoder-decoder (V enc-dec)** in this layer, whose source is from the **final output of the Encoder stage**.

# The Encoder-Decoder Attention Layer

We directly utilize **the final embedding vectors** generated at the end of the Encoder stage, and multiply those with weight matrices to get **K enc-dec & V enc-dec**. These get used as K and V in this Encoder-Decoder Attention Layer.

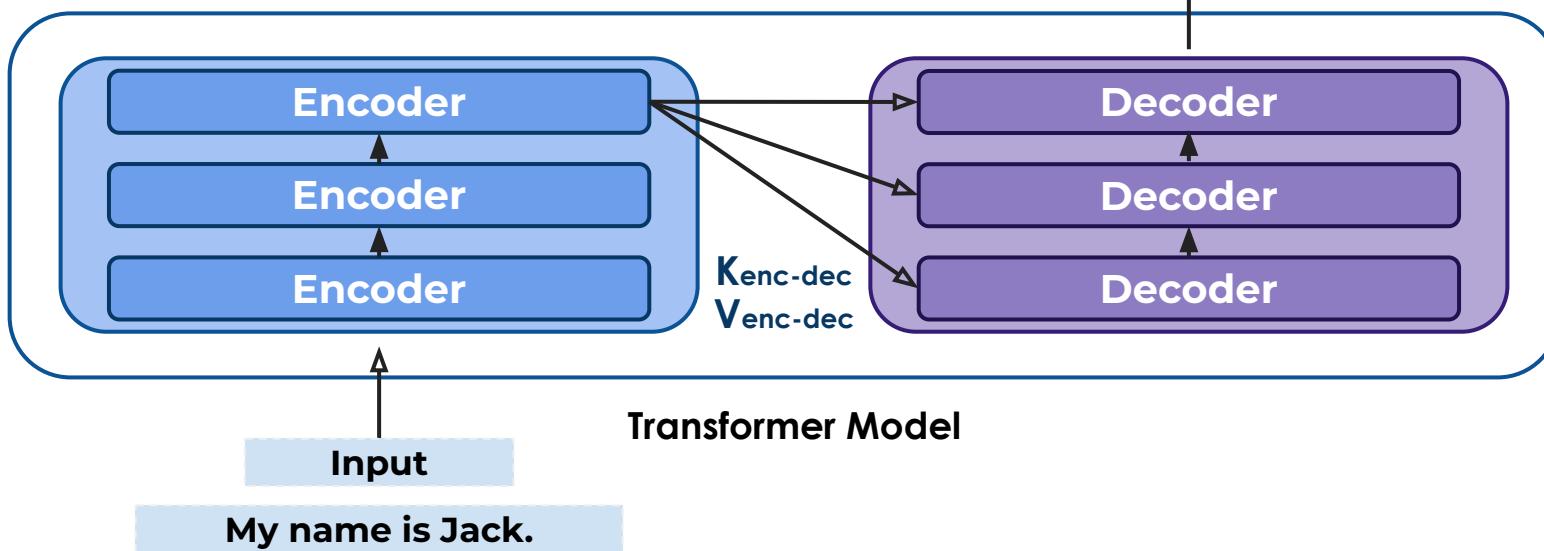
It is also important to mention, that the **Q for < SOS >** (Dec Pos 0) for example, **only relies on the K enc-dec & V enc-dec of the word “I”** (Enc Pos 1) from the input, to predict the word “Ich”. This happens for every Decoder word.

**It is only the Q vector that this layer creates from the input to it**, the way that normally happens in the Self-Attention Layer (where all three of K, Q & V are directly created from the input embeddings to the layer).



# The Encoder-Decoder Attention Layer

This is why the Encoder-Decoder architecture is actually represented as the **final Encoder block feeding every block in the Decoder stage.**



The arrows from the final Encoder block to each Decoder block represent the **K<sub>enc-dec</sub> & V<sub>enc-dec</sub>** from the final Encoder layer being used in the Encoder-Decoder Attention Layer of each Decoder block in the Decoder stage.

# Transformers Quiz

Which of the following best describes BERT's architecture ?

A

BERT is a decoder-only model

B

BERT is an encoder-only model

C

BERT consists of both encoder and decoder layers

# Transformers Quiz

Which of the following best describes BERT's architecture ?

A

BERT is a decoder-only model

B

BERT is an encoder-only model

C

BERT consists of both encoder and decoder layers

# Different Transformer Architectures

There are broadly three-types of transformer models today, based on their usage of Encoder and Decoder blocks

**Encoder-Decoder**

**Encoder-only**

**Decoder-only**

# Encoder-Decoder Transformer

There are broadly three-types of transformer models today, based on their usage of Encoder and Decoder blocks

**Encoder-Decoder**

Utilize the Encoder and Decoder blocks in tandem, similar to the original transformer architecture

Typically used in tasks where the output heavily relies on the input, like **Machine Translation** and **Text Summarization**

Examples: **T5** and **FLAN-T5**

**Encoder-only**

**Decoder-only**

# Encoder-only Transformer

There are broadly three-types of transformer models today, based on their usage of Encoder and Decoder blocks

Encoder-Decoder

Encoder-only

Decoder-only

Utilize only Encoder blocks to generate continuous embeddings from the input

Typically used in discriminative tasks that require embeddings, like for **Text Classification** and **Semantic Search**

Examples: **BERT** and **DistilBERT**

# Decoder-only Transformer

There are broadly three-types of transformer models today, based on their usage of Encoder and Decoder blocks

**Encoder-Decoder**

**Encoder-only**

**Decoder-only**

Utilize only Decoder blocks to auto-regressively predict\* the next token based on the input

Typically used in generative tasks like **Sentence Completion** and **Question-Answering**

Examples: **GPT** and **Llama**

\* Autoregressive prediction involves **Predicting future values based on past values** shared by [priyadarsh@gmail.com](mailto:priyadarsh@gmail.com) only.



# Happy Learning !



This file is meant for personal use by erikvdesigner@gmail.com only.  
Sharing or publishing the contents in part or full is liable for legal action.